

# Optimizer Fusion: Efficient Training with Better Locality and Parallelism

Zixuan Jiang, Jiaqi Gu, Mingjie Liu, Keren Zhu, David Z. Pan

Electrical and Computer Engineering, The University of Texas at Austin

## Machine learning frameworks




- Machine learning algorithms and frameworks **coevolve**.
- Static (symbolic) and dynamic (eager, imperative) computation flow.
- Two critical components.

Automatic differentiation

Iterative optimization methods

## Automatic differentiation



- Compute gradients based on chain rule.

## Iterative optimization methods

- Optimization algorithms in general form.

**Input:** objective function  $f$   
Initialize the starting point  $\theta^{(0)}$   
**for**  $t = 1, 2, \dots$  **do**  
    **if** stopping criterion is met **then**  
        **return**  $\theta^{(t-1)}$   
    **end if**  
     $\Delta\theta = \pi(f, \theta^{(0)}, \theta^{(1)}, \dots, \theta^{(t-1)})$   
     $\theta^{(t)} = \theta^{(t-1)} + \Delta\theta$   
**end for**

- Gradient descent

$$\pi_1 = -\eta \nabla f(\theta^{(t-1)})$$

- Gradient descent with momentum

$$\pi_2 = -\eta \sum_{\tau=0}^{t-1} \alpha^{t-\tau-1} \nabla f(\theta^{(\tau)})$$

- Newton's method

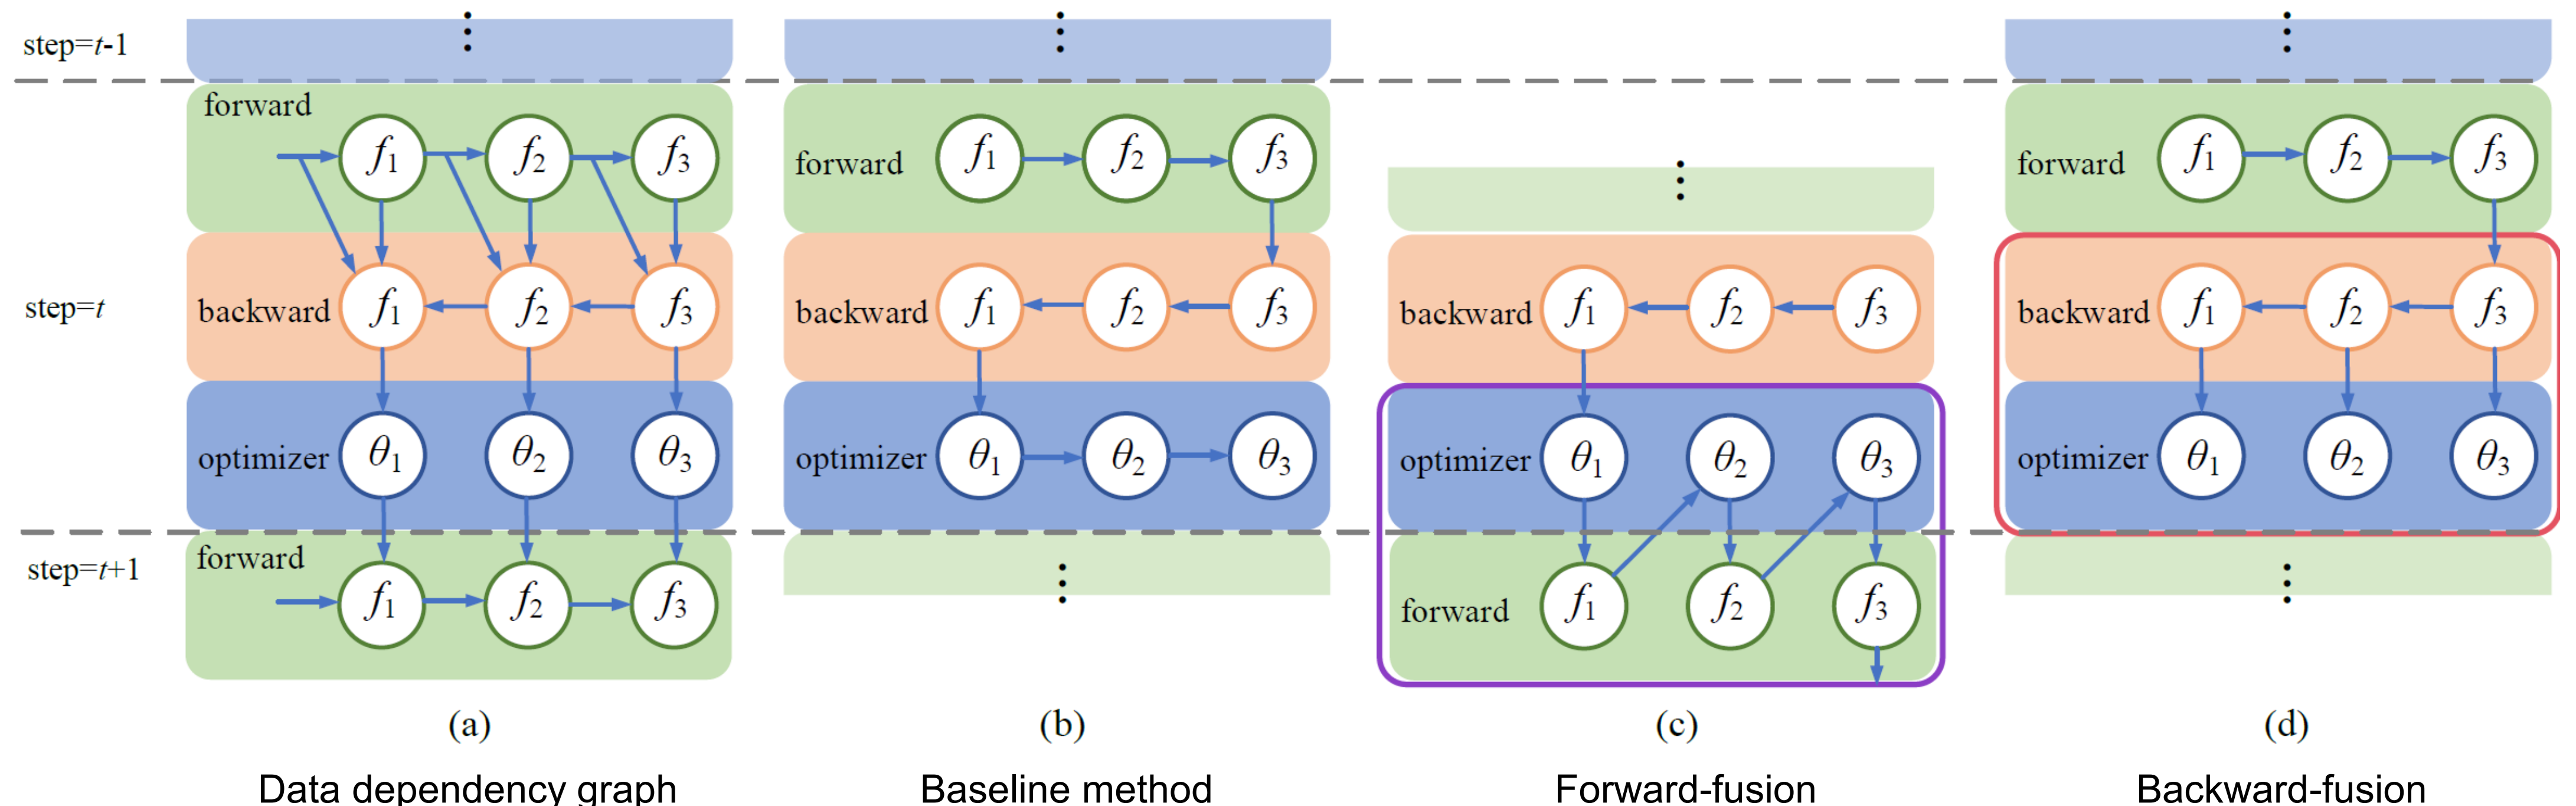
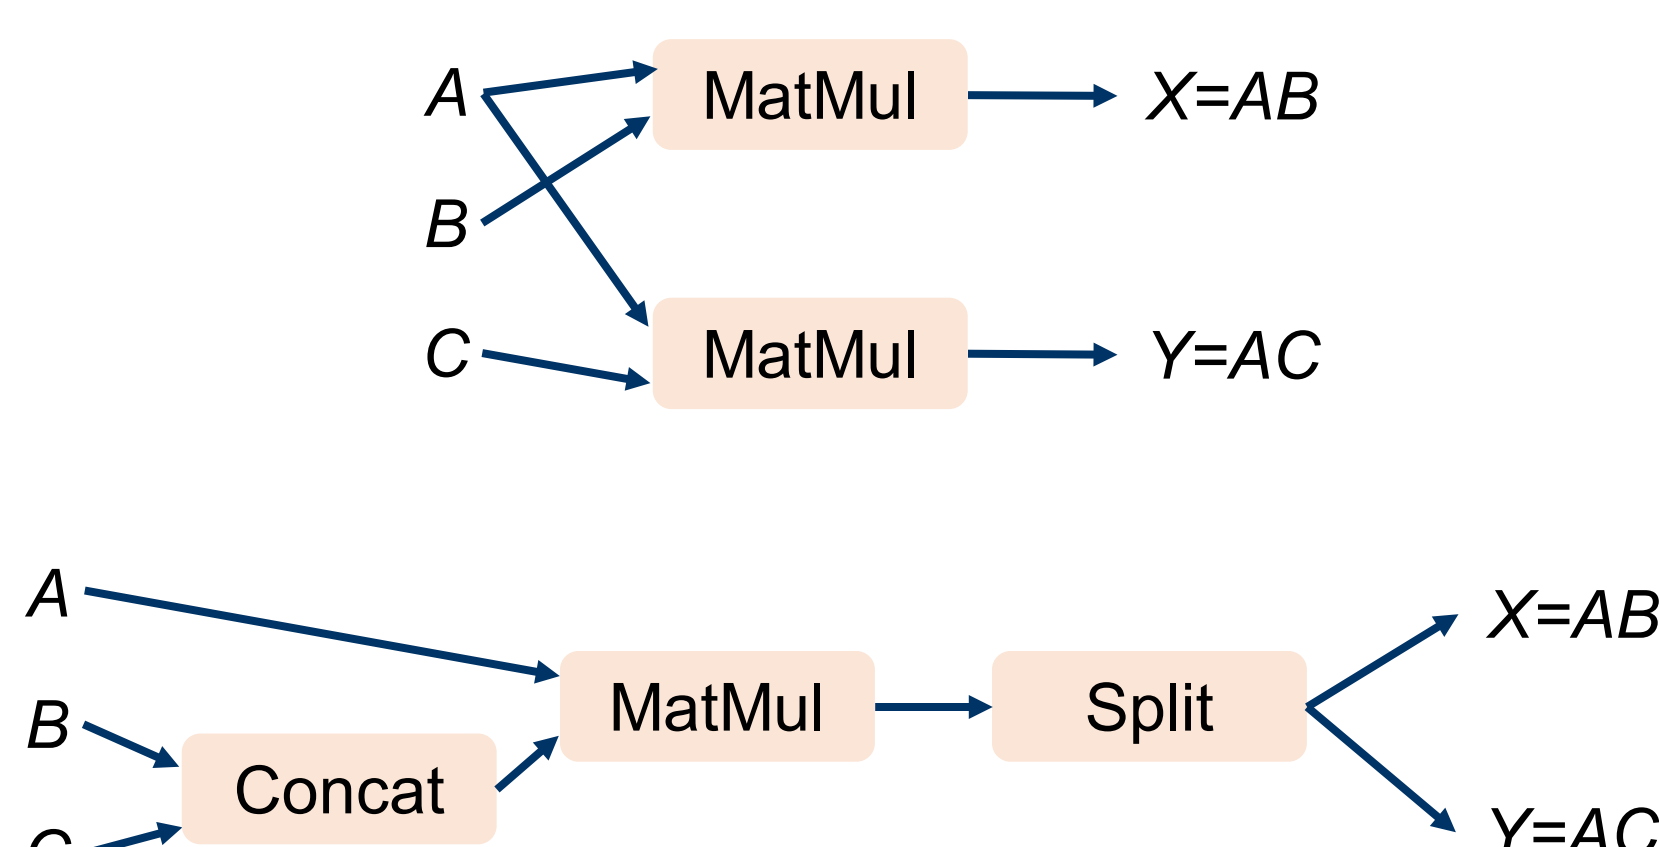
$$\pi_3 = -\eta \nabla^2 f(\theta^{(t-1)})^{-1} \nabla f(\theta^{(t-1)})$$

## Graph optimization


- Operator fusion  
Fuse convolution, batch normalization, and ReLU.




- Operator substitution



## Our methods




## Forward-fusion



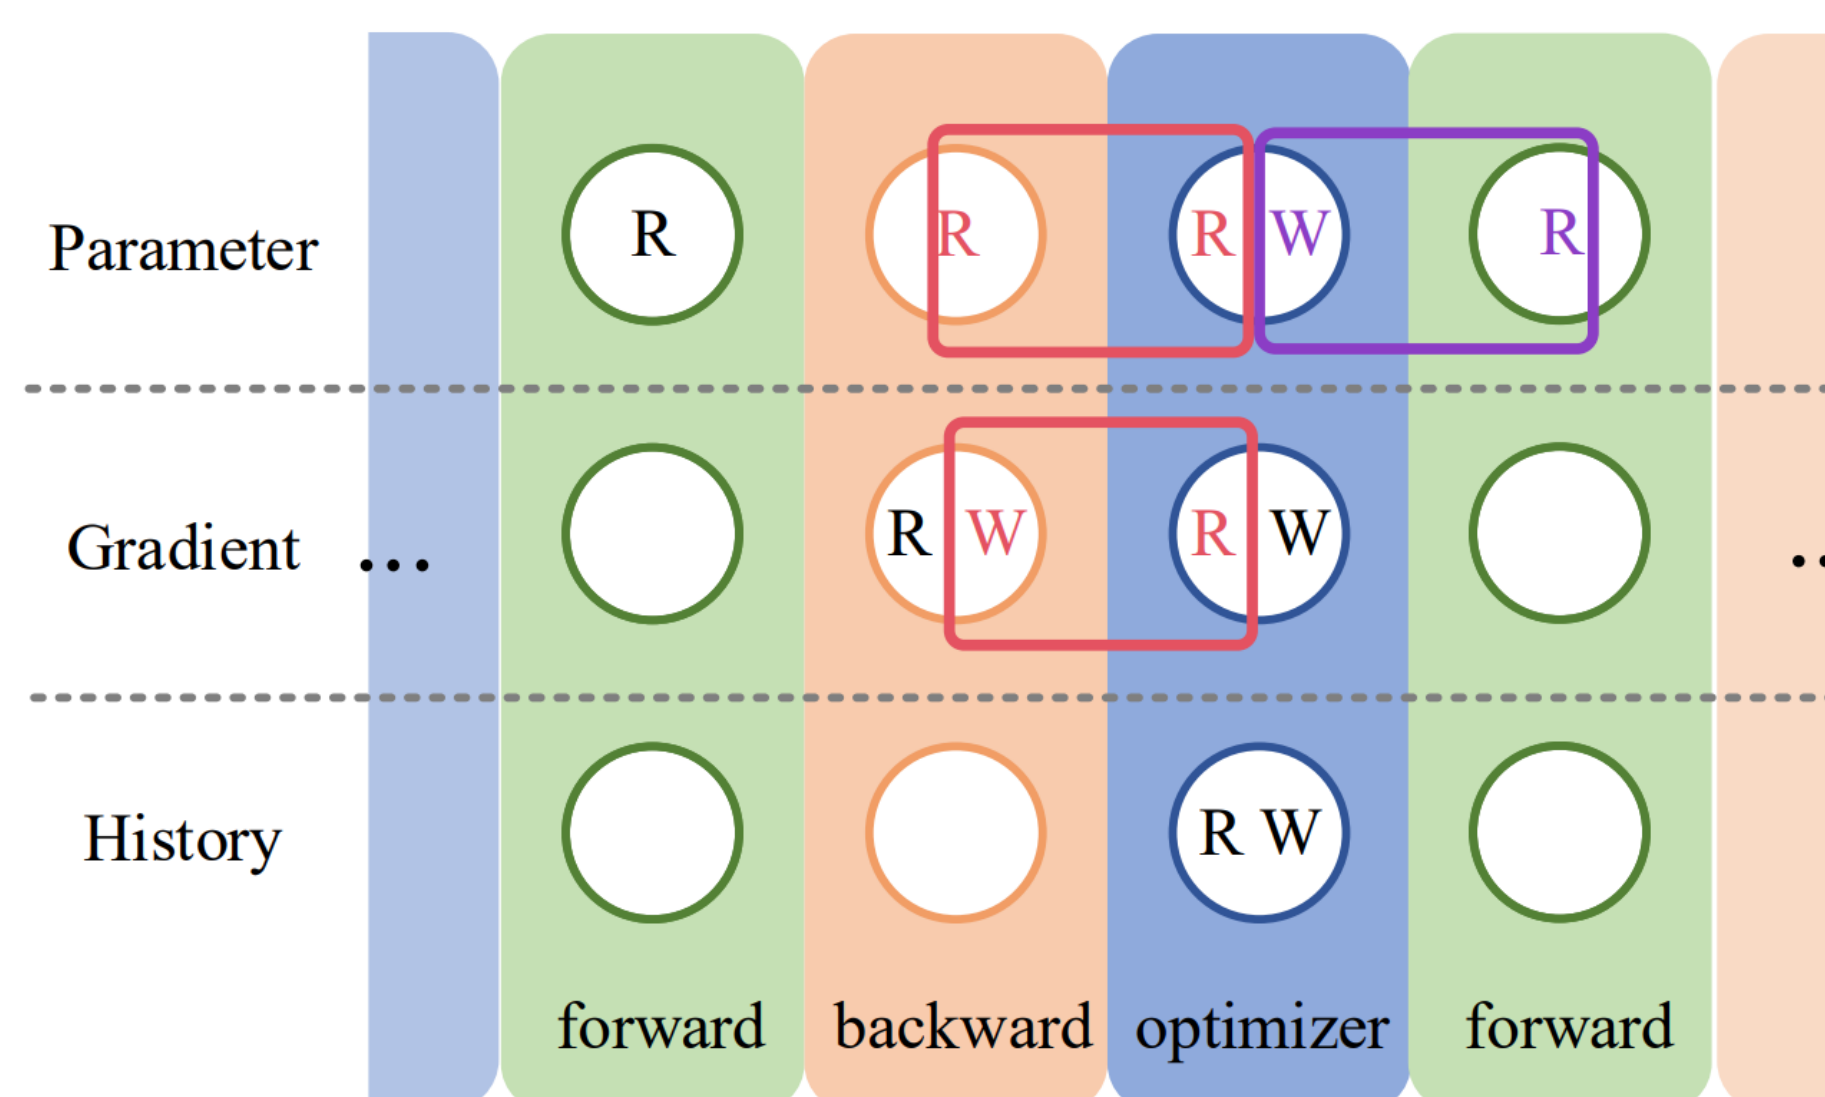
- Fuse the parameter update with the next forward pass.
- Update parameters as **late** as possible.

## Backward-fusion



- Fuse the parameter update with the current backward pass.
- Update parameters as **early** as possible.

## Locality and parallelism




- Memory transactions and data locality in the training process.  $R$  and  $W$  represent memory read and write, respectively. History means parameter history needed in the optimizer, e.g., momentum.

Method	Locality	Parallelism	Global information
Baseline	×	×	✓
Forward-fusion	✓	×	✓
Backward-fusion	✓	✓	×

- In some iterative optimization methods, the gradients will be post-processed after the all the gradients are available. An example is clipping gradients by its norm.
- The *forward-fusion* method is applicable when the global information is needed.

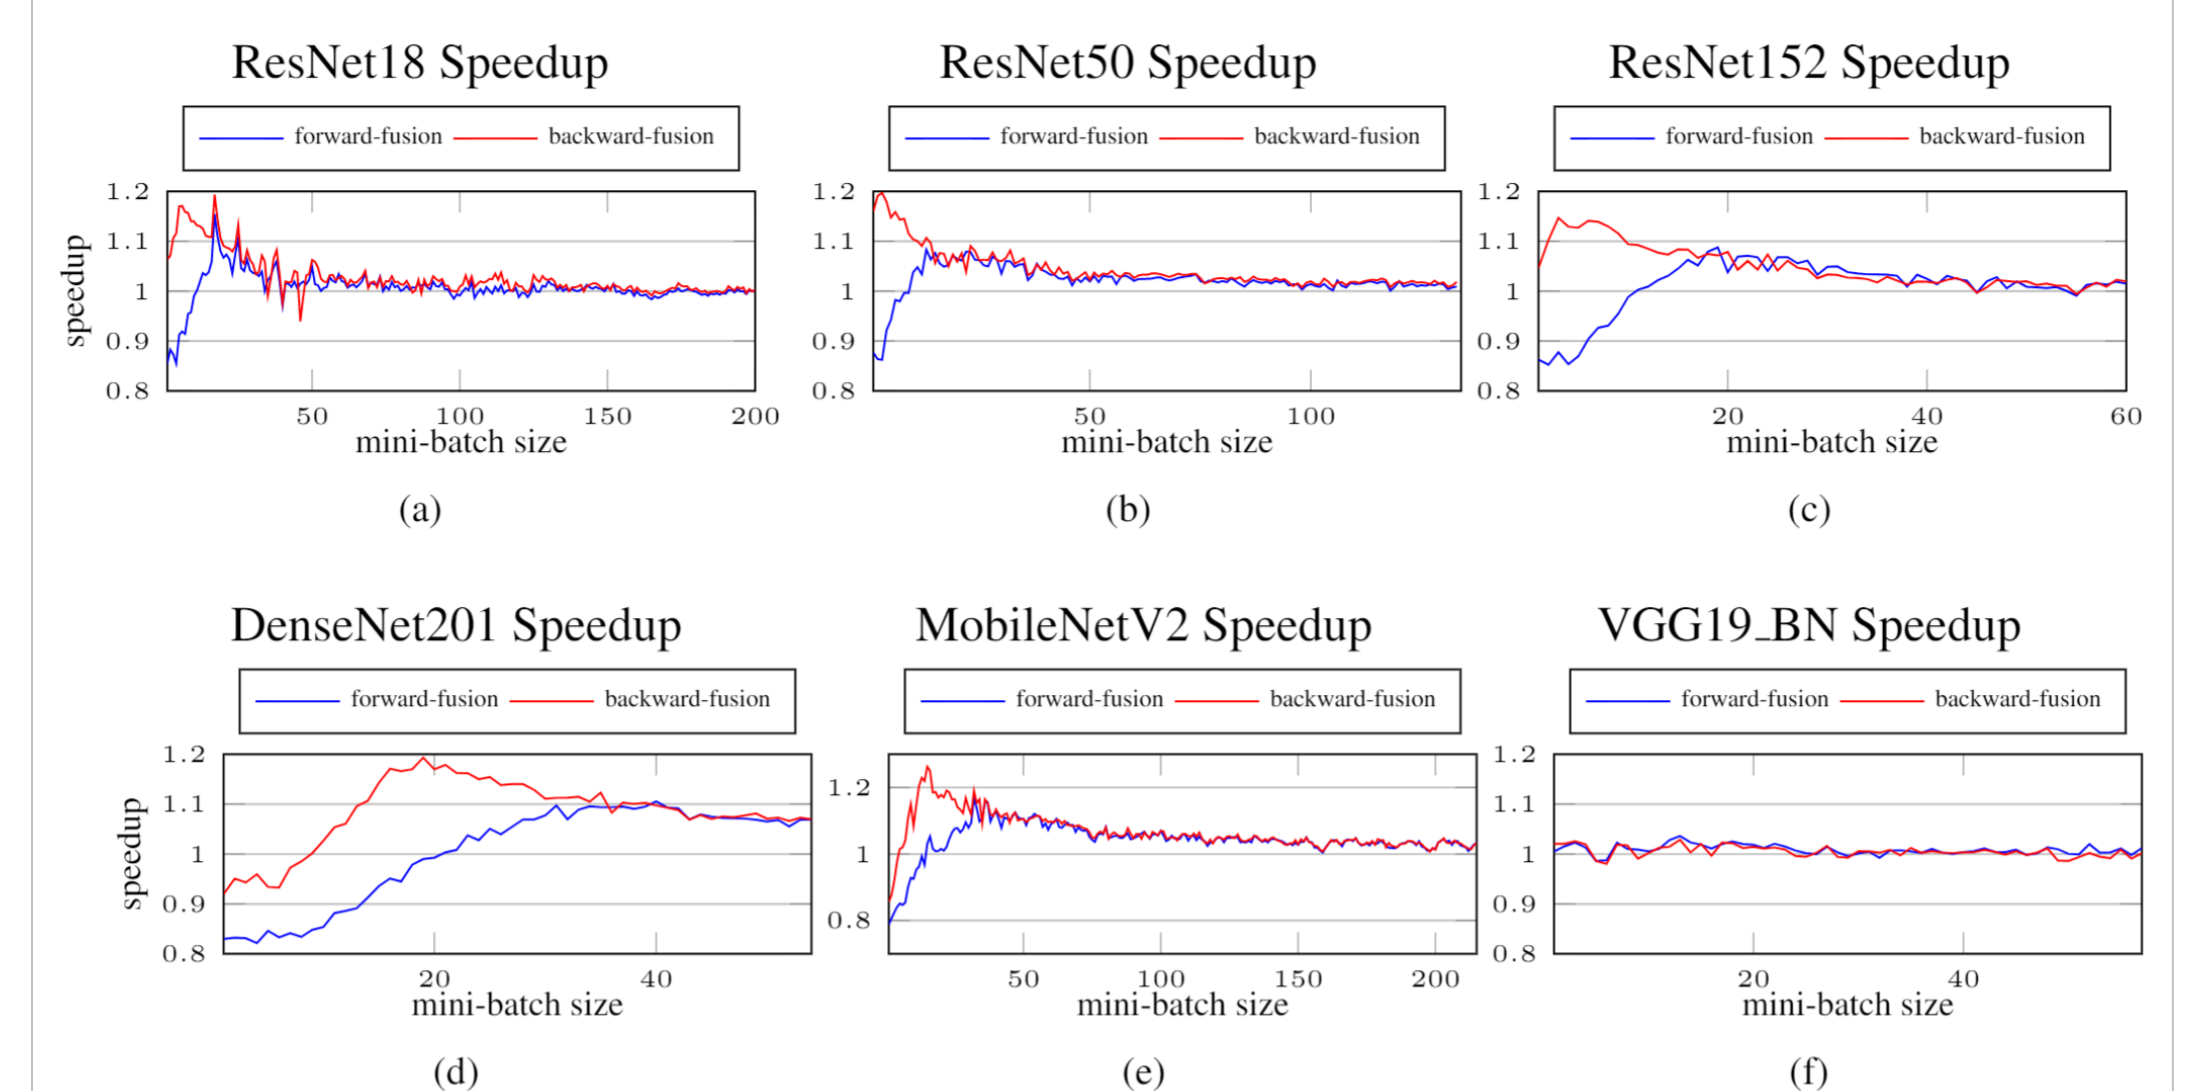
## Experiments



- Training time breakdown of MobileNetV2 with mini-batch size 32. FF, BF are short for *forward-fusion*, *backward-fusion*.
- Our *forward-fusion* and *backward-fusion* improve the training throughput by 12% and 16%.


## Experiments (continued)

### Various mini-batches




- The absolute training time saved by our methods is independent of the mini-batch size.
- The relative speedup will decrease as the mini-batch size grows.

### Various models



- The smaller the average number of parameters per layer, the more locality we can leverage so that our methods can achieve higher training speed.

### Various optimizers



- The horizontal axis represents the ratio of the optimizer time to a whole iteration time.
- The more runtime-costly the optimizer, the higher speedup we can achieve.

## Conclusion

- Conventional eager execution in machine learning frameworks separate the updating of trainable parameters from forward and backward computations.
- We propose two methods forward-fusion and backward-fusion to better leverage the locality and parallelism.
- Experimental results demonstrate the effectiveness and efficiency of our methods across various configurations.
- Our proposed methods are orthogonal to other optimization methods and do not affect the training results. We keep all the features of the eager execution.